

Discrete Data Problems

1 Discrete Random Variable with Two Categories: Binomial Model

1.1 Bernoulli trials and the binomial statistic

Many statistical problems involve repeated experiments with only two possible outcomes:

success/failure, yes/no, event/no event.

The fundamental model for such data is the Bernoulli distribution.

Let $X_1, \dots, X_n$ be independent Bernoulli(p) trials, where

$$\mathbb{P}(X_i = 1) = p, \quad \mathbb{P}(X_i = 0) = 1 - p.$$

Thus, X_i is an indicator of success on trial i .

Mean and variance. A basic calculation gives

$$\mathbb{E}(X_i) = p, \quad \text{Var}(X_i) = p(1 - p).$$

The variance is largest at $p = 1/2$, reflecting maximum uncertainty.

Binomial count of successes. The total number of successes is

$$B = \sum_{i=1}^n X_i.$$

Since the trials are independent and identically distributed,

$$B \sim \text{Binomial}(n, p),$$

with probability mass function

$$\mathbb{P}(B = b) = \binom{n}{b} p^b (1 - p)^{n-b}, \quad b = 0, 1, \dots, n.$$

Interpretation. The binomial model is the discrete analogue of many “counting” procedures in nonparametrics, including the sign test, where we count positive versus negative signs.

Likelihood viewpoint. The binomial likelihood for p is

$$L(p | b) = \binom{n}{b} p^b (1-p)^{n-b},$$

and the maximum likelihood estimator is

$$\hat{p} = \frac{B}{n}.$$

1.2 Exact binomial test

Suppose we want to test whether the success probability equals some benchmark value p_0 :

$$H_0 : p = p_0 \quad \text{vs} \quad H_A : p \neq p_0,$$

with one-sided alternatives defined similarly.

Under H_0 , the distribution of B is exactly known:

$$B \sim \text{Binomial}(n, p_0).$$

Exact p-value. If we observe $B = b_{\text{obs}}$, then:

- For $H_A : p > p_0$,

$$p\text{-value} = \mathbb{P}(B \geq b_{\text{obs}}).$$

- For $H_A : p < p_0$,

$$p\text{-value} = \mathbb{P}(B \leq b_{\text{obs}}).$$

- For $H_A : p \neq p_0$, the p-value is based on both tails, typically doubling the smaller tail probability.

Rejection regions (exact). Let $b_{\alpha;n,p_0}$ denote an upper α critical value:

- For $H_A : p > p_0$, reject if $B \geq b_{\alpha;n,p_0}$.
- For $H_A : p < p_0$, reject if $B \leq c_{\alpha;n,p_0}$.
- For $H_A : p \neq p_0$, reject in both tails:

$$B \geq b_{\alpha_1;n,p_0} \quad \text{or} \quad B \leq c_{\alpha_2;n,p_0}, \quad \alpha_1 + \alpha_2 = \alpha.$$

Remark 1.1 (Discrete nature of exact tests). Because B takes only integer values, it is not always possible to construct a test with size exactly equal to a chosen α . Instead, the achieved significance level may be slightly smaller than α . This is why exact tests can be conservative.

1.3 Large-sample approximation for the binomial test

When n is large, the binomial distribution is well approximated by a normal distribution.

Mean and variance under H_0 .

$$\mathbb{E}(B) = np_0, \quad \text{Var}(B) = np_0(1 - p_0).$$

Define the standardized statistic

$$Z = \frac{B - np_0}{\sqrt{np_0(1 - p_0)}}.$$

Central Limit Theorem. Under H_0 , as $n \rightarrow \infty$,

$$Z \stackrel{\text{approx}}{\sim} N(0, 1).$$

Equivalently,

$$Z^2 \approx \chi_1^2,$$

which connects the binomial test to Pearson chi-square methods.

Rejection regions (normal approximation).

- $H_A : p > p_0$: reject if $Z \geq z_\alpha$.
- $H_A : p < p_0$: reject if $Z \leq -z_\alpha$.
- $H_A : p \neq p_0$: reject if $|Z| \geq z_{\alpha/2}$.

Continuity correction. Since B is discrete but the normal distribution is continuous, a continuity correction sometimes improves accuracy:

$$Z = \frac{B + 0.5 - np_0}{\sqrt{np_0(1 - p_0)}}.$$

Remark 1.2 (When is the approximation valid?). The normal approximation works best when both expected counts are not too small:

$$np_0 \gtrsim 5, \quad n(1 - p_0) \gtrsim 5.$$

If p_0 is very close to 0 or 1, exact methods are preferred.

1.4 Simulation illustrations

Simulation helps visualize the transition from discrete binomial to normal shape.

Small sample ($n = 10$).

```
n = 10; p0 = 1/2; nsim = 1000
B = rbinom(nsim, size = n, prob = p0)

hist(B, breaks = 0:10,
     main = "Binomial (n=10, p=0.5)",
     xlab = "Number of successes")

qqnorm(B); qqline(B)
```

Larger sample ($n = 100$).

```
n = 100; p0 = 1/2; nsim = 10000
B = rbinom(nsim, size = n, prob = p0)

hist(B, breaks = 30,
      main="Binomial(n=100, p=0.5)",
      xlab="Number of successes")

qqnorm(B); qqline(B)
```

As n increases, the histogram becomes smoother and more bell-shaped, illustrating the CLT in action.

2 Example: Sensory Difference Tests (Triangle Test)

2.1 Design

Triangle tests are widely used in food science and sensory evaluation.

- Each of n panelists receives three samples in random order.
- Two samples are identical, one is different.
- The panelist must identify the odd sample.

Under pure guessing, the success probability is

$$p_0 = \frac{1}{3}.$$

We test

$$H_0 : p = \frac{1}{3} \quad \text{vs} \quad H_A : p > \frac{1}{3}.$$

2.2 Large-sample test statistic

Let B be the number of correct identifications. Then under H_0 ,

$$B \sim \text{Binomial}\left(n, \frac{1}{3}\right).$$

The standardized statistic is

$$Z = \frac{B - n(1/3)}{\sqrt{n(1/3)(2/3)}}.$$

Example. With $n = 50$ and observed $B = 25$,

$$Z_{\text{obs}} = \frac{25 - 50(1/3)}{\sqrt{50(1/3)(2/3)}} = 2.5.$$

At $\alpha = 0.05$ one-sided, reject if $Z \geq 1.645$. Thus we reject H_0 .

2.3 p-value: exact vs approximate

Approximate p-value:

$$p_{\text{approx}} = 1 - \Phi(2.5) \approx 0.0062.$$

Exact p-value:

$$p_{\text{exact}} = \mathbb{P}(B \geq 25), \quad B \sim \text{Binomial}(50, 1/3) \approx 0.0108.$$

Remark 2.1. The discrepancy occurs because the normal approximation is imperfect for moderate sample sizes and skewed probabilities like $p_0 = 1/3$. Exact binomial inference remains the gold standard when feasible.

3 Power Calculation (Normal Approximation)

Power is the probability of rejecting H_0 when the alternative H_A is true.

Suppose we test

$$H_0 : p = p_0 \quad \text{vs} \quad H_A : p > p_0,$$

where $B \sim \text{Binomial}(n, p)$.

Using the normal approximation under H_0 , the standardized statistic is

$$Z = \frac{B - np_0}{\sqrt{np_0(1 - p_0)}}.$$

At significance level $\alpha = 0.05$ (one-sided),

$$z_{0.05} = 1.645.$$

Thus we reject H_0 when

$$Z \geq 1.645.$$

This is equivalent to rejecting when

$$B \geq np_0 + 1.645\sqrt{np_0(1 - p_0)}.$$

Power is computed under the true value p (for example $p = 0.6$):

$$\text{Power}(p) = \mathbb{P}_p(\text{Reject } H_0) = \mathbb{P}_p\left(B \geq np_0 + 1.645\sqrt{np_0(1 - p_0)}\right).$$

Under the alternative,

$$B \sim \text{Binomial}(n, p),$$

and we again approximate using a normal distribution:

$$B \approx N(np, np(1 - p)).$$

We standardize using the *true* mean and variance:

$$\mathbb{P}_p\left(\frac{B - np}{\sqrt{np(1 - p)}} \geq \frac{np_0 + 1.645\sqrt{np_0(1 - p_0)} - np}{\sqrt{np(1 - p)}}\right).$$

Since

$$\frac{B - np}{\sqrt{np(1-p)}} \approx N(0, 1),$$

the power is approximately

$$\text{Power}(p) \approx 1 - \Phi\left(\frac{np_0 + 1.645\sqrt{np_0(1-p_0)} - np}{\sqrt{np(1-p)}}\right).$$

This expression is the approximate *power function* of the test.

Example

If $p_0 = 0.5$, $p = 0.6$, and $n = 100$,

$$\text{Cutoff} = 50 + 1.645\sqrt{25} = 50 + 8.225 = 58.225.$$

Under $p = 0.6$,

$$\mu = 60, \quad \sigma = \sqrt{100(0.6)(0.4)} = \sqrt{24} \approx 4.90.$$

Thus

$$\text{Power} \approx \mathbb{P}\left(Z \geq \frac{58.225 - 60}{4.90}\right) = \mathbb{P}(Z \geq -0.36) = \Phi(0.36) \approx 0.64.$$

So the test has roughly 64% power at $p = 0.6$ with $n = 100$.

4 Estimator and Confidence Interval for a Success Probability

4.1 Point estimator

The natural estimator is

$$\hat{p} = \frac{B}{n}.$$

4.2 Standard error

$$\text{se}(\hat{p}) = \sqrt{\frac{p(1-p)}{n}}, \quad \widehat{\text{se}}(\hat{p}) = \sqrt{\frac{\hat{p}(1-\hat{p})}{n}}.$$

4.3 Wald confidence interval

$$\hat{p} \pm z_{\alpha/2} \sqrt{\frac{\hat{p}(1-\hat{p})}{n}}.$$

Remark 4.1 (Limitations). The Wald CI performs poorly when n is small or $\hat{p}$ is near 0 or 1. Exact or Wilson intervals are often preferred.

5 Discrete Random Variable with More Than Two Categories

5.1 Pearson Chi-Squared Goodness-of-Fit Test

Let X take k categories with probabilities $p_1, \dots, p_k$.

Observed counts:

$$X_1, \dots, X_k, \quad \sum X_i = n.$$

Test

$$H_0 : p_i = p_i^0.$$

Pearson statistic:

$$\chi^2 = \sum_{i=1}^k \frac{(X_i - np_i^0)^2}{np_i^0}.$$

Under H_0 ,

$$\chi^2 \approx \chi_{k-1}^2.$$

Why $k - 1$ degrees of freedom? Because probabilities satisfy $\sum p_i = 1$, one category is determined by the others.

Validity condition. Expected counts must satisfy

$$np_i^0 \gtrsim 5.$$

6 Example: Mendel's Pea Plant Experiments (3:1 Ratio)

Mendel's pea plant experiments provide one of the most classical and historically important examples of a *goodness-of-fit* (GOF) test.

Mendel studied traits such as flower color, seed shape, and pod color. For each trait, the offspring were classified into two categories:

- **Dominant phenotype** (e.g., purple flowers),
- **Recessive phenotype** (e.g., white flowers).

Mendelian prediction: the 3:1 ratio

Under simple Mendelian inheritance (one gene with two alleles), the expected probabilities are

$$p_{\text{dom}} = \frac{3}{4} = 0.75, \quad p_{\text{rec}} = \frac{1}{4} = 0.25.$$

Suppose for a given trait we observe counts

$$O_{\text{dom}}, \quad O_{\text{rec}}, \quad \text{with total } n = O_{\text{dom}} + O_{\text{rec}}.$$

Under the null hypothesis H_0 (Mendel's genetic model), the *expected* counts are

$$E_{\text{dom}} = 0.75n, \quad E_{\text{rec}} = 0.25n.$$

Chi-square goodness-of-fit test

To measure how close the observed counts are to the expected Mendelian ratio, we compute the chi-square GOF statistic:

$$\chi^2 = \frac{(O_{\text{dom}} - E_{\text{dom}})^2}{E_{\text{dom}}} + \frac{(O_{\text{rec}} - E_{\text{rec}})^2}{E_{\text{rec}}}.$$

Because there are two categories and no parameters are estimated from the data, this statistic has approximately a chi-square distribution with

$$\text{df} = 2 - 1 = 1.$$

Thus, under H_0 ,

$$\chi^2 \sim \chi_1^2.$$

Combining evidence across multiple traits

Mendel reported results across several traits (for example, 7 traits). For each trait j , one can compute a separate statistic:

$$\chi_{1,j}^2.$$

If the traits are approximately independent, then the sum satisfies

$$\sum_{j=1}^7 \chi_{1,j}^2 \sim \chi_7^2.$$

This provides an overall test of whether Mendel's results agree with the theoretical Mendelian ratios across all traits.

When is the chi-square approximation valid?

The chi-square approximation works well when expected counts are not too small. A common rule of thumb is:

$$E_{\text{dom}}, E_{\text{rec}} \geq 5.$$

If expected counts are very small, exact binomial tests or simulation-based methods may be preferred.

Interpretation

- A small χ^2 value indicates close agreement with Mendel's prediction.
- A large χ^2 value suggests deviation from the 3:1 ratio, possibly due to chance, experimental issues, or model failure.

This example highlights how chi-square GOF tests provide a quantitative way to compare observed categorical data with theoretical probability models.

7 More Than One Discrete Random Variable: Contingency Tables

So far, we have focused on goodness-of-fit tests for a *single* categorical random variable.

In many applications, however, we observe *two* discrete random variables simultaneously, and we want to study whether they are related.

This leads to the framework of *contingency tables*.

Two-way tables

A contingency table records counts for two categorical variables:

- X takes values in categories $i = 1, \dots, r$,
- Y takes values in categories $j = 1, \dots, c$.

The observed count in cell (i, j) is

$$O_{ij} = \#\{(X, Y) \in (i, j)\}.$$

All counts together form an $r \times c$ table:

	Y_1	Y_2	$\dots$	Row total
X_1	O_{11}	O_{12}	$\dots$	$n_{1\cdot}$
X_2	O_{21}	O_{22}	$\dots$	$n_{2\cdot}$
$\vdots$	$\vdots$	$\vdots$	$\ddots$	$\vdots$
X_r	O_{r1}	O_{r2}	$\dots$	$n_{r\cdot}$
Column total	$n_{\cdot 1}$	$n_{\cdot 2}$	$\dots$	n

Marginal totals

The row totals are

$$n_{i\cdot} = \sum_{j=1}^c O_{ij},$$

and the column totals are

$$n_{\cdot j} = \sum_{i=1}^r O_{ij}.$$

The grand total sample size is

$$n = \sum_{i=1}^r \sum_{j=1}^c O_{ij}.$$

Testing independence

A central question in contingency table analysis is:

Are X and Y independent?

The null hypothesis is

$H_0 : X \text{ and } Y \text{ are independent.}$

Under independence, the joint probability factors:

$$P(X = i, Y = j) = P(X = i) P(Y = j).$$

This implies expected counts of the form

$$E_{ij} = \frac{n_{i\cdot} n_{\cdot j}}{n}.$$

Chi-square test for contingency tables

To compare observed and expected counts, we compute the chi-square statistic:

$$\chi^2 = \sum_{i=1}^r \sum_{j=1}^c \frac{(O_{ij} - E_{ij})^2}{E_{ij}}.$$

Under H_0 , this statistic is approximately chi-square distributed with

$$\text{df} = (r - 1)(c - 1).$$

Interpretation

- If χ^2 is small, the data are consistent with independence.
- If χ^2 is large, there is evidence that X and Y are associated.

Contingency table methods generalize goodness-of-fit ideas to the study of relationships between multiple categorical variables.

8 Testing Two Success Probabilities: 2×2 Tables

A very common problem in biostatistics and clinical research is comparing two success probabilities:

- group 1 has success probability p_1 ,
- group 2 has success probability p_2 .

Examples include:

- treatment vs placebo response rates,
- disease prevalence in exposed vs unexposed groups,
- survival vs death proportions in two populations.

The 2×2 table setup

Data are summarized in a 2×2 contingency table:

	Success	Failure	Total
Group 1	X_1	$n_1 - X_1$	n_1
Group 2	X_2	$n_2 - X_2$	n_2
Total	$X_1 + X_2$	$n - (X_1 + X_2)$	n

The sample proportions are

$$\hat{p}_1 = \frac{X_1}{n_1}, \quad \hat{p}_2 = \frac{X_2}{n_2}.$$

The main hypothesis of interest is

$$H_0 : p_1 = p_2 \quad \text{vs} \quad H_a : p_1 \neq p_2.$$

8.1 Fisher's exact test

When sample sizes are small, the usual chi-square approximation may not be accurate, especially if some expected cell counts are less than 5.

In this setting, Fisher's exact test provides an *exact* test of association.

Key idea: conditioning on margins

Fisher's test conditions on the observed row and column totals. That is, we treat the margins as fixed and ask:

Given these margins, how unusual is the observed table under H_0 ?

Under H_0 (equal success probabilities), the number of successes in group 1, X_1 , follows a hypergeometric distribution:

$$X_1 \sim \text{Hypergeometric}(n_1, X_1 + X_2, n).$$

The probability of observing a particular table is

$$P(X_1 = x) = \frac{\binom{X_1 + X_2}{x} \binom{n - (X_1 + X_2)}{n_1 - x}}{\binom{n}{n_1}}.$$

Exact p-value

The p-value is computed by summing probabilities of all tables that are as extreme or more extreme than the observed one.

- Fisher's test is valid for any sample size.
- It is especially recommended when expected counts are small.

When is Fisher's test used?

- Small clinical trials,
- Rare outcomes,
- Sparse contingency tables.

8.2 Large-sample z-test for $p_1 - p_2$

When sample sizes are sufficiently large, we can use a normal approximation to compare two proportions.

The difference in sample proportions is

$$\hat{p}_1 - \hat{p}_2.$$

Pooled estimator under the null

Under the null hypothesis $H_0 : p_1 = p_2 = p$, we estimate the common probability by pooling both groups:

$$\hat{p}_d = \frac{X_1 + X_2}{n_1 + n_2}.$$

Test statistic

The large-sample z-statistic is

$$Z = \frac{\hat{p}_1 - \hat{p}_2}{\sqrt{\hat{p}_d(1 - \hat{p}_d) \left(\frac{1}{n_1} + \frac{1}{n_2} \right)}}.$$

Under H_0 , we have approximately

$$Z \sim N(0, 1).$$

Thus, a two-sided p-value is

$$p\text{-value} = 2P(|Z| > |z_{\text{obs}}|).$$

Connection to chi-square

Squaring the z-statistic gives

$$Z^2 \approx \chi_1^2,$$

so this test is equivalent to the Pearson chi-square test in the 2×2 case.

When is the z-test appropriate?

A common rule of thumb is that all expected counts should satisfy

$$E_{ij} \geq 5.$$

9 2×2 Pearson Chi-Squared Tests

Another standard approach to comparing two proportions is the Pearson chi-square test for independence in a 2×2 contingency table.

Expected counts under independence

Let O_{ij} denote the observed cell counts.

Under the null hypothesis of independence, the expected counts are

$$E_{ij} = \frac{n_{i\cdot} n_{\cdot j}}{n},$$

where

- $n_{i\cdot}$ are the row totals,
- $n_{\cdot j}$ are the column totals,
- n is the grand total.

Pearson chi-square statistic

The Pearson statistic is

$$\chi^2 = \sum_{i=1}^2 \sum_{j=1}^2 \frac{(O_{ij} - E_{ij})^2}{E_{ij}}.$$

Under H_0 and for large samples,

$$\chi^2 \approx \chi_1^2.$$

Interpretation

- Small χ^2 means observed counts match independence.
- Large χ^2 provides evidence of association between group and outcome.

Continuity correction (Yates). Because the chi-square distribution is continuous while the data are discrete, the Pearson approximation can slightly overstate significance in small samples.

Yates' continuity correction adjusts for this by subtracting 0.5:

$$\chi_{\text{Yates}}^2 = \sum_{i,j} \frac{(|O_{ij} - E_{ij}| - 0.5)^2}{E_{ij}}.$$

Effect of the correction

- Makes the test more conservative (larger p-values),
- Mainly used when counts are small,
- Often replaced today by Fisher's exact test in sparse tables.

10 Measures and Visualizations of Association

The chi-square test tells us whether there is evidence against independence, but it does not explain *where* the association comes from.

In practice, after rejecting H_0 , we want to answer:

Which cells contribute most to the departure from independence?

This motivates residual-based diagnostics and visualization tools.

10.1 Residual-based interpretation

Recall that the Pearson chi-square statistic is

$$\chi^2 = \sum_{i,j} \frac{(O_{ij} - E_{ij})^2}{E_{ij}}.$$

This shows that the overall test statistic is built from cell-by-cell contributions.

Pearson residuals

A convenient standardized measure of deviation in each cell is the *Pearson residual*:

$$r_{ij} = \frac{O_{ij} - E_{ij}}{\sqrt{E_{ij}}}.$$

These residuals compare the observed count to the expected count under independence, scaled by the standard deviation $\sqrt{E_{ij}}$.

Interpretation

- If $r_{ij} \approx 0$, then the cell count is close to what independence predicts.
- If $r_{ij} > 0$, then

$$O_{ij} > E_{ij},$$

meaning there are *more observations than expected* in that category combination.

- If $r_{ij} < 0$, then

$$O_{ij} < E_{ij},$$

meaning there are *fewer observations than expected*.

Which cells drive the association?

Cells with large absolute residuals,

$$|r_{ij}| \gg 0,$$

are the main contributors to the chi-square statistic and therefore drive the rejection of independence.

As a rough guideline:

- $|r_{ij}| > 2$ suggests a notable departure,

- $|r_{ij}| > 3$ suggests a strong departure.

Thus, residuals provide a diagnostic tool beyond the single p-value.

10.2 Association plots

While residuals can be examined numerically, it is often more informative to visualize them.

Association plots provide a graphical representation of the deviations between observed and expected cell counts.

Main idea

Under the null hypothesis of independence:

$$O_{ij} \approx E_{ij}, \quad r_{ij} \approx 0,$$

so the plot should show bars close to zero.

When independence fails, some cells will show large positive or negative residuals, highlighting the source of association.

R code example

In R, the function `assocplot()` produces an association plot:

```
assocplot(x,
          xlab="death penalty",
          ylab="gun registration",
          main="Association plot")
```

Here, `x` is a contingency table object.

How to read an association plot

- Bars above zero indicate **excess counts**:

$$O_{ij} > E_{ij}.$$

- Bars below zero indicate **deficits**:

$$O_{ij} < E_{ij}.$$

- The height of the bar is proportional to the magnitude of the residual.
- Large residual patterns reveal which category combinations are most strongly associated.

Connection to hypothesis testing

- Under the null hypothesis of independence, residuals should fluctuate randomly around zero.
- In real data, large systematic residuals indicate departures from independence, leading to rejection of H_0 .
- Association plots help interpret the chi-square test by showing *where* the dependence occurs.

10.3 Mosaic plots as an alternative visualization

Another popular visualization for contingency tables is the *mosaic plot*.

- The area of each tile represents the observed frequency.
- Deviations from independence can be highlighted by shading.
- Mosaic plots are especially useful for larger tables.

Thus, graphical tools such as association plots and mosaic plots provide important insight beyond the overall chi-square test result.